

# Kubernetes Tutorial

Kubernetes (or k8s for short) is a container orchestration system that allows us to deploy Docker containers across a cluster of machines. It automates scheduling and assignment of containers to physical machines, as well as fault-tolerance, scaling, and application lifecycle management. This tutorial gives a quick-start introduction to using Kubernetes with our testing cluster, but you will have to refer to additional material to get the full picture. We try to recommend useful references throughout.

## Prerequisites

You should have the following setup already completed before you begin:

1. A username and password for our Gitlab instance at [git.webis.de](https://git.webis.de)
2. If you are located **outside Weimar university**, you should have obtained a VPN certificate from [cert-server.webis.de](https://cert-server.webis.de), and have your OpenVPN client configured correctly to use it. If you have never worked with us before, the person who introduced you will have set this up for you.

## Connecting to the Test Cluster

### Install the *kubectl* commandline client

Refer to the official tutorial on [kubernetes.io](https://kubernetes.io), and complete the appropriate “Install kubectl on \$OPERATING\_SYSTEM” step. For subsequent configuration, return here and proceed below.

### Authenticate to the cluster and Configure *kubectl*

Visit [openid.webis.de/k8s-login](https://openid.webis.de/k8s-login) to authenticate to the Kubernetes cluster with your Gitlab credentials. After successful authentication, you will see a sequence of commands that you should paste into your terminal. Do so.

### Test the Connection

Run the following command in your terminal:

```
1 kubectl get nodes
```

You should see output similar to the following:

1	NAME	STATUS	ROLES	AGE	VERSION
2	webis6	Ready	<none>	132d	v1.14.1
3	webis7	Ready	<none>	132d	v1.14.1
4	webis8	Ready	<none>	132d	v1.14.1
5	webis9	Ready	<none>	132d	v1.14.1
6	webis10	Ready	<none>	132d	v1.14.1

## Create a Personal Namespace

Namespaces isolate different applications on a Kubernetes cluster. For now, we'll create one for you to do the initial tutorial steps in. In the following, we assume that [username] is a name that uniquely identifies you — substitute something appropriate, preferably your university login or similar.

```
1 kubectl create namespace [username]
2 kubectl config set-context webis6 --namespace=[username]
```

## Kubernetes Architecture

A Kubernetes cluster consists of several daemon processes. We will review only the two most important ones below. You can find much more in-depth information in the [relevant documentation](#).

### kube-apiserver

Runs on a single physical node that serves as the cluster manager (although failover is possible). Handles all communication of the k8s cluster with the outside world (especially with clients that want to deploy applications), as well as within the cluster (for example, nodes report their state and resource availability)

The apiserver for our testing cluster runs on the machine `webis6.medien.uni-weimar.de`.

### kubelet

Runs on every cluster node where containers will be deployed. Works with the local `docker` daemon to achieve this. Receives instructions from the apiserver, and communicates its status there.

## A Brief Tour of Some Important Concepts in Kubernetes

Below, we will very briefly review some of the basic concepts you may encounter when deploying your code on Kubernetes.

### General Procedure

The `kubectl` client on your local machine talks to the cluster's `kube-apiserver` — they communicate using the **Kubernetes REST API** (if you really wanted to, you could use something like `curl` as a client instead). Whenever you want anything to happen on the cluster, you send a declarative specification of the desired state, expressed in terms of the basic resource objects specified by the API. In what follows, we'll use **YAML** to express our resource specifications.

For each of the examples below, we'll follow this basic procedure:

1. Create a text file containing the resource description, say `my-resource.yaml`.
2. Send it to the cluster using the `kubectl` client. For example:

```
bash kubectl apply -f my-resource.yaml
```

*(instead of `apply`, other operations, such as `patch`, `replace`, or `delete` are available)*

We'll walk through doing this for the most important types of resources.

### Pods

**Pods** are the smallest deployable units in Kubernetes. A pod consists of one or more Docker containers that together perform a single task.

### A First Pod

Create a file `my-first-pod.yaml` with the following contents:

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: my-first-pod
5 spec:
6   containers:
7     - name: my-container
8       image: alpine
9       command: ['sh', '-c', 'echo Hello, Kubernetes!']
```

This file describes your pod, in particular, it specifies which containers it consists of, and which commands they should run. In this case, we're spawning a single container using the standard alpine image (a minimal Linux distribution). The container's root process will spawn a shell that prints a simple message to stdout and then terminates immediately.

Deploy this pod like this:

```
1 kubectl apply -f my-first-pod.yaml
```

Then look at its state and output:

```
1 kubectl get pod
2 kubectl describe pod my-first-pod
3 kubectl logs my-first-pod
```

By the time you get to the `kubectl get`, it will likely have already terminated. The cluster will, by default, restart it a few times until eventually giving up (Pods are expected to run indefinitely).

We can get rid of it now:

```
1 kubectl delete pod my-first-pod
```

## A Second Pod

Let's define a new pod that doesn't immediately terminate:

`my-second-pod.yaml`

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: my-second-pod
5 spec:
6   containers:
7     - name: my-container
8       image: alpine
9       command: ['sh', '-c', 'echo Hello, Kubernetes! &&
      ↪ touch /keep-running && while [ -e /keep-
      ↪ running ]; do sleep 10; done && echo Bye!']
```

The command has become a bit more complex: the container now creates an empty file on startup, and as long as that file still exists it will keep running.

We'll send our pod specification to the cluster:

```
1 kubectl apply -f my-second-pod.yaml
```

In the list of running pods (`kubectl get pod`), the Kubernetes cluster will eventually show our new pod as running, and it will stay that way. We can run additional commands in an existing pod, which is very useful for debugging, because you can spawn an interactive shell:

```
1 kubectl exec -it my-second-pod sh
```

*(the option `-i` means you want to pass stdin to the pod, and `-t` that stdin is a terminal)*

Inside that new shell, we can write a new file, and then make it terminate:

```
1 echo "hello" > /my-new-file
2 cat /my-new-file
3 rm /keep-running
```

The pod will terminate within ten seconds and your shell session along with it, but the cluster will restart the pod immediately. However, the file you created is gone:

```
1 kubectl exec my-second-pod -- cat /my-new-file

1 cat: can't open '/my-new-file': No such file or directory
2 command terminated with exit code 1
```

## Volumes

All storage in the pod we've just created is ephemeral, that means any on-disk changes are lost once it terminates or restarts. **Volumes** are Kubernetes' storage abstraction, and they supply pods with persistent state.

Kubernetes supports many different storage backends, but we'll only look at a small number of examples below.

### Short-term Persistence with *emptyDir*

We'll create a new pod with a modified spec.

`pod-with-emptydir.yaml`

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-with-emptydir
5 spec:
6   volumes:
7     - name: my-volume
8       emptyDir: {}
```

```

9   containers:
10  - name: my-container
11    image: alpine
12    command: ['sh', '-c', 'echo Hello, Kubernetes! &&
      ↳ touch /keep-running && while [ -e /keep-
      ↳ running ]; do sleep 10; done && echo Bye!']
13    volumeMounts:
14  - mountPath: /myvol
15    name: my-volume

```

Run: `kubectl apply -f pod-with-emptydir.yaml`

We've added two new things here:

1. A new key `volumes` in our pod spec, which contains a list of all the volumes we use, in particular their name, which storage backend they use, and any options (in this case, none). Our volume is named `my-volume` and uses Kubernetes' `emptyDir` storage backend — this creates an empty directory on the Kubernetes host, and shares it with our pod.
2. A new key `volumeMounts` inside our pod's container specification, which specifies which volumes the container should mount were. We could also specify mount options here, e.g. for read-only access.

**Exercise:** Using what you've learned before, spawn a shell inside your pod, create a new file under the `/myvol` directory, and then cause the pod to terminate (by deleting `/keep-running`). After Kubernetes restarts your pod, see if the file is still there.

*Spoiler:* it is.

The contents of `emptyDir` volumes persist across pod restarts. This is possible because whenever a pod terminates due to some error, Kubernetes will reschedule it on the same host. However, once you manually delete your pod, your `emptyDir` volumes are irrevocably gone. This is useful e.g. for caching, where your pods may create temporary data that are useful to keep around, but not critically important.

## PersistentVolumes and *rbd*

**Persistent volumes** exist independently of the pod lifecycle. They supply your Kubernetes pods with storage that persists across multiple deployments, or that can be attached to different pods over time. They do introduce a lot of additional complexity, but we will skip over most of that here (refer to the official documentation if open questions remain).

There are two relevant Kubernetes API objects, **PersistentVolume** and **PersistentVolumeClaim**. For the purposes of this tutorial, you will only interact with the `PersistentVolumeClaim` object, which describes the storage you need.

The cluster will automatically manage the underlying volume that satisfies your request.

`my-first-pvc.yaml`

```
1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: my-first-pvc
5 spec:
6   accessModes:
7     - ReadWriteOnce
8   resources:
9     requests:
10      storage: 1Mi
11   storageClassName: ceph-rbd
```

Run: `kubectl apply -f my-first-pvc.yaml`

The keys under `spec` describe the volume that we want. Here, we’re requesting one megabyte worth of storage, using a `StorageClass` we’ve named `ceph-rbd`. Storage classes are defined by the administrators of a particular Kubernetes cluster to make different storage backends available to Kubernetes — in our case, we’re using a `Ceph` cluster that provides `virtual block devices` to our containers.

You will be able to see your claim in the output of `kubectl get pvc`. The `STATUS` column will show *Pending* as the cluster provisions your volume, and switch to *Bound* once it’s ready for use. Next, we’ll define a Pod that mounts this storage into one of its containers.

Run: `kubectl apply -f pod-with-pv.yaml`

This pod works very similarly to the one with the `emptyDir` before. However, we now use our persistent virtual block device that we’ve just created instead of a temporary directory on the host file system, and mount that in the `/myvol` directory.

**Exercise:** Deploy the pod above to the cluster, spawn a shell inside its container, and write something to a file in the `/myvol` directory. Inspect the output of `kubectl describe pod pod-with-pv` — The *From* column in table at the end of the output tells you on which physical host the pod is running (They are named “webis6” through “webis10”). Delete your pod, and re-deploy it. Repeat until it gets deployed on a different host than it was before. You will find that the contents of `/myvol` are preserved across deployments.

## Cleaning up

Let’s get rid of the pod we just created, and allow the cluster to free up the storage we used:

```
1 kubectl delete pod/pod-with-pv pvc/my-first-pvc
```

## ConfigMaps

**ConfigMaps** can share textual data that is specified directly in the YAML file with your containers. As the name implies, they are intended for configuration, but are also useful to store simple scripts inline, in the same place as your deployment files.

Let's create a ConfigMap, and a Pod that uses it in one go.

**pod-with-configmap.yaml**

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: my-first-configmap
5 data:
6   my-script: |-
7     #!/bin/sh
8     echo "I am a container startup script defined in a
9       ↪ ConfigMap!"
10    echo "I will do nothing for the next hour."
11    echo "By the way: the value of the FOO environment
12       ↪ variable is '${FOO}'."
13    sleep 3600
14    echo "Byeeee"
15   my-other-key: "my other value"
16 ---
17 apiVersion: v1
18 kind: Pod
19 metadata:
20   name: pod-with-configmap
21 spec:
22   volumes:
23     - name: my-configmap-volume
24       configMap:
25         name: my-first-configmap
26         defaultMode: 0700
27   containers:
28     - name: my-container
29       image: alpine
30       command: ['/bin/startup.sh']
31       volumeMounts:
32         - name: my-configmap-volume
33           subPath: my-script
```



```

32     mountPath: /bin/startup.sh
33     readOnly: true
34   env:
35     - name: FOO
36       valueFrom:
37         configMapKeyRef:
38           name: my-first-configmap
39           key: my-other-key

```

Run: `kubectl apply -f pod-with-configmap.yaml`

The example above introduces several new concepts: You can specify multiple API objects in a single YAML file; they are separated by lines containing three dashes. Here, we specify a **ConfigMap object** that contains two keys (the value of **one of them** happens to be the code of a shell script, written with the YAML syntax for multi-line strings). Below the `---`, we define a **pod** that makes use of this ConfigMap, and we see two ways that ConfigMaps can be used:

1. **Mounted into the container's file system.** To this end, we give the pod a **volume** of the configMap type that refers to our ConfigMap above (we also **specify** that its contents should be mounted with execute permissions). In the container's volumeMounts key, we then **mount** the value of the ConfigMap's my-script key as the file /bin/startup.sh (we introduce some new options to volume mounts that allow us to **mount an individual key as a file in an existing directory**). We could also mount the entire ConfigMap in an empty directory as we did before, and would then get one file per key). We use /bin/startup.sh as **the command** the container should execute.
2. **Via environment variables.** In the env key of the container specification, we **set** the value of the environment variable FOO to the value of our ConfigMap's second key. Note that the startup script **prints out** the value of this environment variable.

Use `kubectl logs pod-with-configmap` to see the output of the startup script, and verify that this works.

## Pod Networking

According to the **Kubernetes networking model**, each pod in the cluster gets its own IP address, in a virtual, cluster-internal IP range. Every pod can communicate with every other pod given its IP, no matter where it runs in the cluster (containers inside the same pod share an IP address and can communicate via the localhost interface<sup>\*\*</sup>). We will quickly demonstrate this.

**Exercise:** Start any two of the pods we've used in this tutorial before (except for the first one that terminates immediately). Using `kubectl` you can see their

internal IP, and which physical host they have been deployed on (we'll use a custom output format to show the columns we're interested in):

```
1 kubectl get pod -o custom-columns=Name:metadata.name,Host
  ↳ :spec.nodeName,PodIP:status.podIP
```

The output will look something like this:

1 Name	Host	PodIP
2 my-second-pod	webis10	10.23.151.179
3 pod-with-configmap	webis6	10.23.204.126

Our two pods run on different physical hosts in this case (you may delete and re-create one of them if that's not the case for you). Open two terminals side-by-side, and spawn a shell in each of the two pods.

In the first pod, run

```
1 nc -v -l -p 1234
```

*(This starts **netcat** in listening mode on Port 1234)*

Switch to the terminal with the shell for the other pod, and run

```
1 nc -v <ip-of-the-first-pod> 1234
```

*(Refer to the table you produced with `kubectl get` above for the correct IP)*

Type any line of text followed by enter, it should show up in the other pod's terminal (this also works in the other direction). Type Ctrl+C to get out of netcat.

## Deployments

When you run complex applications on a Kubernetes cluster, you will generally not create individual pods directly. Instead, you will create a **Deployment** which declaratively specifies a set of pods that you would like to exist. This has several advantages over creating individual pods: for example, you can easily specify a pod template that should be replicated in a given number of copies for load-balancing and failover purposes. When you change the desired number of replicas, the cluster will create/destroy pods automatically to match.

To illustrate, we will use the [flask server example](#) from our docker tutorial. Let's build this with a tag and push it to Dockerhub:

```
1 cd examples/flask-server
2 docker build . -t <your-dockerhub-username>/flask-server
3 docker push <your-dockerhub-username>/flask-server
```

(You may have to *create an account* and run `docker login first`)

Once it's on Dockerhub, we can deploy it to the kubernetes cluster as well. We will use the aforementioned Deployment controller this time:

`flask-server-deployment.yaml`

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: flask-server-deployment
5   labels:
6     app: my-flask-server
7 spec:
8   replicas: 1
9   selector:
10    matchLabels:
11      app: my-flask-server
12   template:
13     metadata:
14       labels:
15         app: my-flask-server
16     spec:
17       containers:
18       - name: flask-server
19         image: webis/flask-server
```

Run `kubectl apply -f flask-server-deployment.yaml`

When looking at the file, we note a few novelties: First of all, it *specifies* a Deployment API object. The key `spec.template` *contains* the description of the pods that this deployment will create — this subtree has the same keys as the pods that we've created directly up to now. Second, we are using Kubernetes' *Labels and Selectors* mechanism for the first time. In brief, labels allow us to annotate any object we create with arbitrary key-value pairs. We use this to give both *our deployment* and *our pods* a common `app` label. This is necessary, so that the Deployment can identify the pods that it manages — we tell it how to identify its pods using *the selector key*.

When you run `kubectl get all`, you will notice that a total of three new objects have been created. Example output might look like this:

NAME	READY	STATUS	RESTARTS	AGE
pod/flask-server-deployment-6d9c7f59fc-gvc4x	1/1	Running	0	16s
NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/flask-server-deployment	1/1	1	1	16s
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/flask-server-deployment-6d9c7f59fc	1	1	1	16s

Our deployment controller manages a *ReplicaSet*, which it turn manages the pods, and ensures that the desired number are running.

**Exercise:** Change the pod spec to use **your own image**. Increase the **number of replicas** in the deployment spec. Run `kubectl apply` again and observe what happens via `kubectl get all`. *(you should find that the desired number of replicas start with the new image. Only after they are up will your old pod be terminated.)*

## Publishing Ports

Of course, we'll actually want to use the flask server at some point. In order to make that happen, we have to make the port that flask uses known to the cluster. For that, we'll deploy a modified version:

**flask-server-deployment-ports.yaml**

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: flask-server-deployment
5   labels:
6     app: my-flask-server
7 spec:
8   replicas: 1
9   selector:
10    matchLabels:
11      app: my-flask-server
12   template:
13     metadata:
14       labels:
15         app: my-flask-server
16     spec:
17       containers:
18         - name: flask-server
19           image: webis/flask-server
20           ports:
21             - name: flask-server
22               containerPort: 5000
23               protocol: TCP
```

Run: `kubectl apply -f flask-server-deployment-ports.yaml`

The **only addition** is the `ports` key in the container section of the pod specification. It assigns the name `flask-server` to the internal container port 5000.

You can access this port in a local browser by running:

```
1 kubectl port-forward deployment.apps/flask-server-
  ➔ deployment 12345:flask-server
```

Then, visit `localhost:12345` in your browser. You should see the “Hello World” message from the flask app.

## Services

If we want to make our little flask app visible to the world, the approach with `kubectl port-forward` is less than ideal. This is where **Services** come in. While there is a multitude of different implementations (many of them specific to the infrastructure provided by public clouds such as **AWS** or **GCE**), we will only cover the basic **NodePort** services here. These work as follows: you assign a port number to your service, in a range allowed by the cluster; once the service is deployed, *every physical cluster node* listens on this port number, and forwards any connections to it to your service. The following example will illustrate this:

`flask-service.yaml`

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: flask-service
5   labels:
6     app: my-flask-server
7 spec:
8   type: NodePort
9   ports:
10    - port: 80
11      targetPort: flask-server
12      nodePort: 31999
13   selector:
14     app: my-flask-server
```

Run: `kubectl apply -f flask-service.yaml`

Since NodePorts must be unique cluster-wide, you will have to change **this port number** until you find one that nobody else is using; as currently configured, our cluster allows ports from 30000 up to 32767 for NodePort services. Find one that’s available.

Just like Deployments, services identify their associated pods with **a selector**. Every service has **a name** and (at least) one **service port** that **is mapped** to the named port from our deployment. The service’s name functions as an internal DNS name: in our case, you can access it as `http://flask-service:80` from anywhere inside the kubernetes cluster. To illustrate, spawn a shell in one of our previous pods, for example:

```
1 kubectl exec -it my-second-pod sh
```

Once in, install curl:

```
1 apk add --no-cache curl
```

And use it to access the flask app:

```
1 curl flask-service:80
```

You should see the Hello world page's HTML source.

In addition we have defined the public-facing NodePort. Assuming the number you used above is 31999, you should be able to access the service in your browser using any of the following URLs:

```
1 http://webis6:31999/
2 http://webis7:31999/
3 http://webis8:31999/
4 http://webis9:31999/
5 http://webis10:31999/
```

(of course, the service and the deployment can be specified in the same file using the three-dashes syntax)

## Things We Didn't Cover

Kubernetes defines (many) more kinds of API objects. A few that may be relevant are briefly mentioned below:

- **InitContainers** can be added to a pod to run initialization code before the pod's regular containers are created.
- **Secrets** work similarly to ConfigMaps, but are intended for sensitive data like API keys or passwords. They are stored in encrypted form.
- **StatefulSets** manage groups of pods that are created from the same template, but have some run-time state that makes them unique; as an example, consider a distributed database, where each pod runs the same code, but has its own unique database shard.
- **DaemonSets** manage groups of pods where every node (or some subset) of the cluster should run exactly one copy.
- **Jobs** and **CronJobs** define pods that do not run indefinitely as daemon process, but rather are expected to terminate at some point. Jobs run once, CronJobs run periodically.